



Bagging & Boosting

Qinliang Su (苏勤亮)

Sun Yat-sen University

suqliang@mail.sysu.edu.cn

Ensemble Methods Overview

- It is difficult to learn a *strong classifier* that can always classify instances correctly
- But it is easy to learn a lot of *'weak' classifiers*

A weak classifier may not perform well on the whole dataset, but may perform well on a fraction of samples, e.g., some may be good at recognizing 'cat', while some others may be good at recognizing 'dog'
- If weak classifiers perform well on different fractions of samples, it is possible to derive *a strong classifier by combining these weak classifiers in an appropriate way*
- Two questions
 - 1) How to produce these weak classifiers?
 - 2) How to combine the weak classifiers?

Two Types of Combining Methods

1) Unweighted average

- Majority vote

2) Weighted average

- Give better classifiers bigger weighting

For example, consider a two-class classification problem $\{-1, 1\}$

Two basic classifiers: $\hat{y}_1 = \text{sign}(f_1(\mathbf{x}))$ $\hat{y}_2 = \text{sign}(f_2(\mathbf{x}))$

Final classifiers: $\hat{y}_e = \text{sign}(\alpha_1 f_1(\mathbf{x}) + \alpha_2 f_2(\mathbf{x}))$

Remark: The weak classifiers could be of any kind, e.g. decision trees, SVM, neural networks, logistic regression etc.

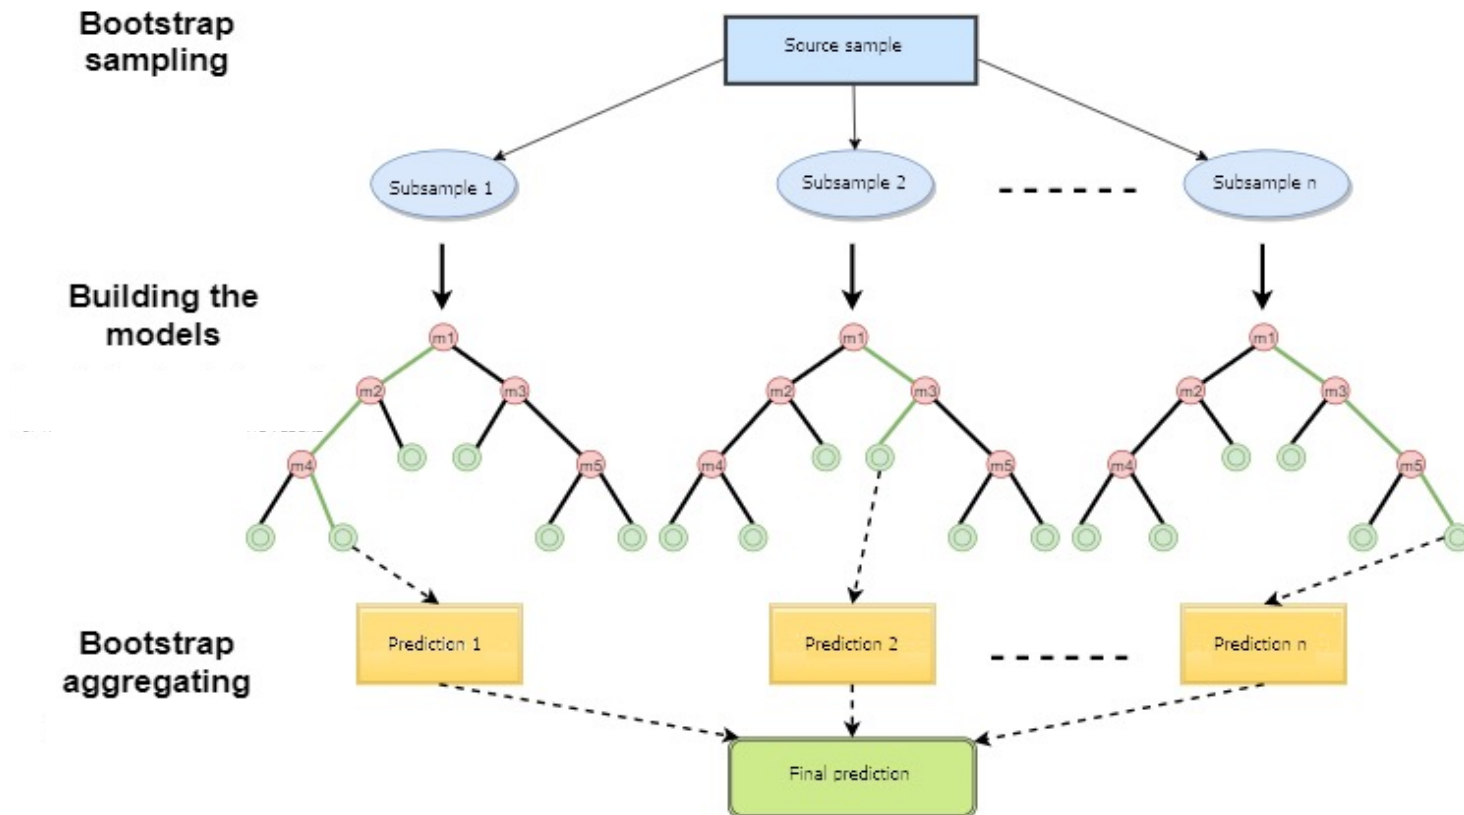
Outline

- Ensemble methods
 - Bagging (majority vote)
 - Boosting (weighted average)

Deriving Weak Classifiers

- We don't know how to obtain classifiers that perform well on different fractions of samples
- Instead, we attempt to obtain classifiers that produce independent prediction errors, that is, encouraging their predictions *to be uncorrelated*. For example,
 - 1) Creating subsets of the training dataset by bootstrapping
 - Randomly draw N' samples from the N -sample training dataset *with replacement*
 - Repeat the above procedure K times, generating subsets S_1, S_2, \dots, S_K
 - 2) Training a decision tree on each of the subset S_k


3) Combine K decision trees into one by majority voting




At testing, pass test data through all K classifiers, using the majority voting result as the final prediction

Example: Bagged Decision Trees


1 tree




Average of 5 trees



Average of 25 trees



Average of 100 trees

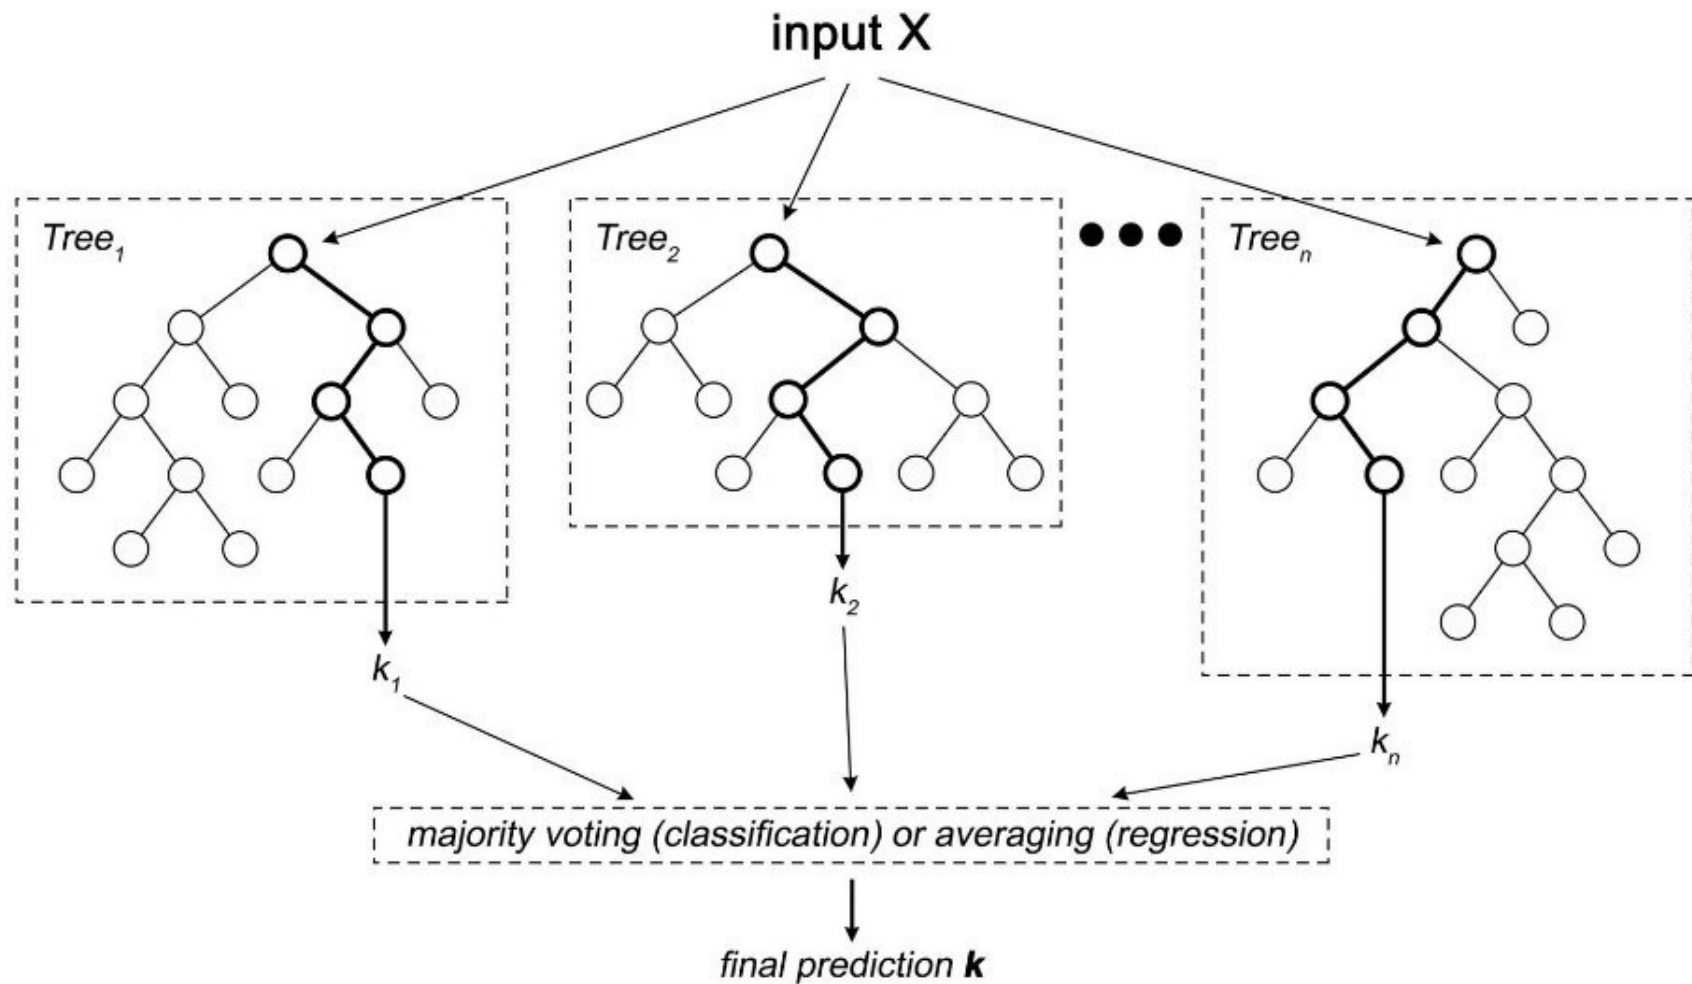


Random Forest

- What happens if the prediction errors in the classifiers derived above are *still strongly correlated* ?
 - ⇒ Combining them by majority vote may not help too much on the final performance
- *Remedy*: Introducing extra randomness into the learning process of decision trees

As building the decision trees, only use a subset of randomly selected attributes

- Random forest illustration



Outline

- Ensemble methods
 - Bagging (majority vote)
 - Boosting (weighted average)

Overview of Boosting Methods

- *Weak classifiers derivation*

Repeat the following steps several times

- 1) Identifying the examples that are incorrectly classified
- 2) Re-training the classifier by *giving more weighting to the misclassified examples*


- *Combining*

Combining the prediction results of each classifier by *weighted average*

How to *weight the examples and prediction* results is the key

Adaboost


- Consider a two-class classification problem with 10 training examples



- For simplicity, only consider the weak classifiers whose decision boundaries are parallel to the axes, that is,

$$\hat{y} = \text{sign}(x_1 + b) \quad \text{or} \quad \hat{y} = \text{sign}(x_2 + b)$$

- First iteration




➤ Error rate of the first classifier h_1 : $\epsilon_1 = 0.3$

➤ Weighting of the classifier h_1 : $\alpha_1 = \frac{1}{2} \ln \left(\frac{1-\epsilon_1}{\epsilon_1} \right) = 0.42$

$$\frac{1 - \epsilon_1}{\epsilon_1} = \frac{\text{correct rate}}{\text{error rate}}$$

⇒ The weighting is *positively proportional* to performance of the classifier




- *Misclassified* examples' weights are **amplified** by e^{α_1}

$$e^{\alpha_1} = \sqrt{\frac{1 - \epsilon_1}{\epsilon_1}} = \sqrt{\frac{\text{correct rate}}{\text{error rate}}}$$

- *Correctly classified* examples' weights are **dampened** by $e^{-\alpha_1}$


- Second iteration



- Error rate of the second classifier h_2 : $\epsilon_2 = 0.21$
- Weighting of the classifier h_2 : $\alpha_2 = \frac{1}{2} \ln \left(\frac{1-\epsilon_2}{\epsilon_2} \right) = 0.65$
- Misclassified examples' weights are amplified by $e^{\alpha_2} = \sqrt{\frac{1-\epsilon_2}{\epsilon_2}}$
- Correctly classified examples' weights are dampened by

$$e^{-\alpha_2} = \sqrt{\frac{\epsilon_2}{1-\epsilon_2}}$$


- Third iteration



- Error rate of the second classifier h_3 : $\epsilon_3 = 0.14$
- Weighting of the classifier h_3 : $\alpha_3 = \frac{1}{2} \ln \left(\frac{1-\epsilon_3}{\epsilon_3} \right) = 0.92$
- Stop the iteration

- Final classifier

Combining the three classifiers with a linear combination



- Adaboost algorithm

- 1) Initialize the weight of examples as $\omega_0^{(n)} = \frac{1}{N}$ for $n = 1, \dots, N$
- 2) For the k -th iteration, train a classifier $h_k(\mathbf{x})$ with the training examples weighted by $w_{k-1}^{(n)}$

- 3) Evaluate the weighted classification error

$$\epsilon_k = \frac{\sum_{n=1}^N \omega_{k-1}^{(n)} I(y_i \neq h_k(\mathbf{x}_n))}{\sum_{n=1}^N \omega_{k-1}^{(n)}}$$


- 4) Determine the vote stake of the k -th classifier

$$\alpha_k = \frac{1}{2} \ln \left(\frac{1 - \epsilon_k}{\epsilon_k} \right)$$

- 5) Update the weights of examples as

$$\omega_k^{(n)} = \omega_{k-1}^{(n)} \exp\{-y_i h_k(\mathbf{x}_i) \alpha_k\}$$

- A typical error rate curve as a function of the number of weak classifiers



- Typical weak classifiers
 - Decision trees
 - Logistic regressions
 - Neural networks

Theories behind the Adaboost

- Define the following *exponential loss*

$$\mathcal{L} = \sum_{n=1}^N \exp\{-y^{(n)} h_{combine}(\mathbf{x}^{(n)})\}$$

where $\mathbf{x}^{(n)}$ is the input, $y^{(n)} \in \{-1, 1\}$ is the label; and $h_{combine}(\cdot)$ is the combined classifier

$$h_{combine}(\mathbf{x}) = \alpha_1 h_1(\mathbf{x}) + \cdots + \alpha_K h_K(\mathbf{x})$$

with $h_k(\mathbf{x})$ representing the k -th component classifier, e.g.,
 $h_k(\mathbf{x}) = \text{sign}(\mathbf{w}_k^T \mathbf{x} + b_k)$

- It can be proved that the Adaboost algorithm is *equivalent to minimize the exponential loss* in a sequential fashion